

High-Performance Clifford Layers in C

TEAM 75

Simon Huber, Anirudhh Ramesh, Konstantinos Chasiotis, Mikkeline Elleby



Eidgenössische Technische Hochschule Zürich
Swiss Federal Institute of Technology Zurich

Clifford Algebra

- Operate on multivectors of scalars, vectors, planes, volumes, etc
- We work with basis vector that represent direction in space
- Uses geometric product with a main key rule:

$$e_i \times e_i = g_i$$

- This results in blades – a geometric object
- For a dimension d, we have 2^d blades which we define how they multiply with each other via a multiplication table. For 1d for instance it looks like this:


x	1	e_1
1	1	e_1
e_1	e_1	g_0

Reference Codebase

- Started with given codebase : [GitHub - microsoft/cliffordlayers](#)
- Created wrapper files for each c implementation to import the function into the original codebase
- Translated into c and focused on files for :
 - linear layer
 - multi-vector sigmoid linear units
 - group normalizations
- Extended the pytest for testing our implementation
 - Compared them with the python implementation


Layer: Linear Layer

Batch=1, channels_in = 3, channels_out = 2, dim=1




Layer: Linear Layer

Batch=1, channels_in = 3, channels_out = 2, dim=1



Layer: Clifford Linear Layer

Batch=1, channels_in = 3, channels_out = 2, dim=1




Batch=1, channels_in = 3, channels_out = 2, dim=1

Layer: Linear Layer

■ BASELINE

Flattened input multiplied by kernels.




Improving

Re-use inputs
 $\text{Weight}[\text{in_ch}][\text{out_ch}][\text{blades}]$
 Have to 'random index' weights

Batch=1, channels_in = 3, channels_out = 2, dim=1

Re-use outputs <- probably better!
 $\text{Weight}[\text{out_ch}][\text{in_ch}][\text{blades}]$

In_ch -> contiguous mem accesses




Problem with this? Strided weight accesses (switch between input channels)?


Blade inner performs very poorly?

When in_channels > cache line, we have cache striding => Worse than input re-ordering + non-strided accesses.

Kernel approach



No-Kernel approach




Improving

Re-use inputs
 $\text{Weight}[\text{in_ch}][\text{out_ch}][\text{blades}]$
 Have to 'random index' weights

Batch=1, channels_in = 3, channels_out = 2, dim=1

Re-use outputs <- probably better!
 $\text{Weight}[\text{out_ch}][\text{in_ch}][\text{blades}]$

In_ch -> contiguous mem accesses



Problem with this? Strided weight accesses (switch between input channels)?


Blade inner performs very poorly?

When in_channels > cache line, we have cache striding => Worse than input re-ordering + non-strided accesses.

Improving

Batch=1, channels_in = 3, channels_out = 2, dim=1

So could multiply both together.
 Apply AVX
 Apply ...
 Clean up all the code
 Roofline plots and stuff



Multiply and store the output variable per input channel based on index!

AVX vectorization of this approach

Should scale without strided access issues.


Any interesting AVX optimizations? Channel-dependent however

Baseline

Different input weight per channel
input_f[batch][blades][input]

Different input weight per channel
weights[blades][output][input]

Different input weight per channel
output_f[batch][blades][output]




We'd get striding on the weight matrix (if inputs >> large)

No kernel w/o flattened

Different input weight per channel
input_f[batch][input][blades]

Different input weight per channel
weights_f[output][input][blade]



We'd get striding as we multiply each input. Ideal for input x **input[batch][input][blades]**

■ Optimization:


Cache the kernel (avoid reconstruction)

During kernel construction, pre-compute various multiple computations

AVX-2 vectorizations of the kernel multiplication, vector k-accumulators, FMA

=> ~3x overall scalar speed-up

=> ~8x overall vectorized speed-up



■ Optimization:


Cache the kernel (avoid reconstruction)

During kernel construction, pre-compute various multiple computations

AVX-2 vectorizations of the kernel multiplication, vector k-accumulators, FMA

=> ~3x overall scalar speed-up

=> ~8x overall vectorized speed-up



Layer: Linear Layer

■ Further Optimization ideas (in progress)

Currently, the kernels look as follows:

1D:

[w0 g0*w1]
[w1 w0]



(Simplified 1D kernel?)



2d:

... (similar pattern)

3d:

... (similar pattern)

Could we halve the size of the kernels (e.g. $k_{opt} = [w0 \ g0*w1]$)?


Then, in the multiplications we can do e.g. $x @ k_{opt}$ and $k_{opt} @ x$ to construct the final output.

Another idea: Construct the kernels in a way that allows no flattening?

=> Results in a smaller kernel (i.e. less memory load cycles)

Layer: Linear Layer (Performance)

Clifford Linear Performance
Intel core i5-1035g4 @ 3.70GHz
Performance [FLOPs/cycle] vs. Batch size




Layer: Linear Layer (Roofline)

Intel Core i5-1035G4 @ 3.70GHz


Compiler: gcc (Ubuntu 13.3.0-6ubuntu2~24.04) 13.3.0

Performance (F/C) vs Operational Intensity (F/B)




Layer: Group normalizations

- Normalizes across a group of channels
- Involves:
 - Calculating and subtracting means
 - Calculating covariance matrices
 - Finding maximums
 - Calculating Cholesky decompositions
 - Back-substitutions




Groupnorm Covariance Optimization

- BASELINE: Transpose into temporary array and do standard matrix multiply
- OPTIMIZATION: Vectorized direct calculation with AVX fmadd_ps




Groupnorm Backsubstitution Optimization

- Dependencies across iterations => hard to optimize single Back substitution
- **BASELINE:** One matrix and one vector at once
- **OPTIMIZATION:** One matrix and 8 vectors at once (AVX)
 - Needs memory layout restructuring



Groupnorm Performance Plots

Clifford Group Normalisation
Intel core i5-1035g4 @ 3.70GHz
Performance [FLOPs/cycle] vs. Batch size




Groupnorm Roofline

Intel Core i5-1035G4 @ 3.70GHz

Compiler: gcc (Ubuntu 13.3.0-6ubuntu2~24.04) 13.3.0

Performance (F/C) vs Operational Intensity (F/B)



Multi-vector sigmoid linear units

1. Pre-activation z :

$$z = b_c + \sum_{j=0}^{I-1} W_{c,j} X_{b,c,p,j}$$

2. Non-linearity $gate$:

$$\text{gate} = \sigma(z) = \frac{1}{1 + e^{-z}}$$

3. Output Y :

$$Y_{b,c,p,i} = X_{b,c,p,i} \times \text{gate}$$

■ Depth-wise 1×1 “convolution” → sigmoid → gate each blade component

```
CLASS torch.nn.Conv2d(in_channels, out_channels, kernel_size, stride=1, padding=0,
dilation=1, groups=1, bias=True, padding_mode='zeros', device=None,
dtype=None) [SOURCE]
```

Machine perception notes
helped us solve a bug ;)

“convolution over an input signal composed of several input planes.”
output value of the layer with input size (N, C_{in}, H, W) and output size $(N, C_{\text{out}}, H, W)$ can be precisely described as:

$$\text{out}(N_i, C_{\text{out},j}) = \text{bias}(C_{\text{out},j}) + \sum_{k=0}^{C_{\text{in}}-1} \text{weight}(C_{\text{out},j}, k) * \text{input}(N_i, k)$$

where $*$ is the valid 2D cross-correlation operator, N is a batch size, C denotes a number of channels, H is a height of input planes in pixels, and W is width in pixels.

Most DL frameworks implement a correlation (and call it convolution).

Annoying, but ...


Indeed 😊

It doesn't matter for learning. In DL we learn the kernel weights, we don't care whether we learn a flipped kernel or not.

MV: Inverse mapping

- Full blade space → Stored slots
 - We start with a large space of possible blade indices (`0...all_blades-1`), but only store a small subset (slots) per channel
- Build tiny lookup table:

```
for (int b = 0; b < all_blades; ++b) inv[b] = -1;  
for (int i = 0; i < slots; ++i) inv[input_blades[i]] = i;
```



Go from $O(\text{batch} \cdot \text{channels} \cdot \text{all_blade_space})$
□ to $O(\text{all_blades})$ memory !

Note : `all_blades` = size of the full blade index range

MV: Fast Sigmoid via Lookup


- Per-element $1/(1+\exp(-z))$ costs tens of cycles each
- Dominates our inner loop □ SO EVEN MORE EXPENSIVE !
 - Idea: Sample once by doing the computation
 - Lookup + interpolate:
 - *Clamp z to [-6,6]*

```
for (int i = 0; i < 512; ++i) {  
    float x = -6.0f + 12.0f * i / 511.0f;  
    sigmoid_table[i] = 1.0f / (1.0f + expf(-x));  
}
```

- No `expf()` inside the hot loop—just two loads + three ops
- Still get ~1e-5 accuracy

Layer: Activation (Performance)

Multivector Sigmoid-linear Unit
Intel core i5-1035g4 @ 3.70GHz
Performance [FLOPs/cycle] vs. Batch size




Layer: Activation (Roofline)

Intel Core i5-1035G4 @ 3.70GHz

Compiler: gcc (Ubuntu 13.3.0-6ubuntu2~24.04) 13.3.0

Performance (F/C) vs Operational Intensity (F/B)



Thank you for listening !